

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import pandas as pd
import io
#data = pd.read_csv(io.BytesIO(uploaded['housing.csv']))
data = pd.read_csv('/content/sample_data/housing.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
data.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populatio
0	-122.23	37.88	41.0	880.0	129.0	322
1	-122.22	37.86	21.0	7099.0	1106.0	2401
2	-122.24	37.85	52.0	1467.0	190.0	496
3	-122.25	37.85	52.0	1274.0	235.0	558
4	-122.25	37.85	52.0	1627.0	280.0	565

## Encoding

(1)Label encoder

(2)Onehot encoder

```
from sklearn.preprocessing import LabelEncoder , OneHotEncoder
data['median_house_value'].value_counts()
```

```
500001.0    965
137500.0    122
162500.0    117
```

```

112500.0    103
187500.0     93
...
303200.0     1
307900.0     1
383200.0     1
360800.0     1
405500.0     1
Name: median_house_value, Length: 3842, dtype: int64

```

```

le=LabelEncoder()
data['median_house_value']=le.fit_transform(data['median_house_value'])
data['median_house_value'].value_counts()

```

```

3841     965
959      122
1209     117
710      103
1459      93
...
3172      1
3275      1
3204      1
3091      1
2119      1
Name: median_house_value, Length: 3842, dtype: int64

```

```
le.classes_
```

```
array([ 14999.,  17500.,  22500., ..., 499100., 500000., 500001.])
```

```
data['ocean_proximity'].value_counts()
```

```

<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: ocean_proximity, dtype: int64

```

```

one_hot = OneHotEncoder()
transformed_data = one_hot.fit_transform(data['ocean_proximity'].values.reshape(-1,1)).toarray()
one_hot.categories_

```

```
array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
      dtype=object)
```

```

transformed_data = pd.DataFrame(transformed_data ,
                                columns = ['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
                                index = data.index)
transformed_data.head()

```

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	0.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	1.0	0.0
2	0.0	0.0	0.0	1.0	0.0

```
transformed_data.iloc[90, ]
```

```
<1H OCEAN    0.0
INLAND       0.0
ISLAND       0.0
NEAR BAY     1.0
NEAR OCEAN   0.0
Name: 90, dtype: float64
```

```
data['median_house_value'][90]
```

```
1209
```

## Normalization & Standardization

```
# consider only numerical columns
```

```
numeric_columns = [c for c in data.columns if data[c].dtype != np.dtype('O')]
numeric_columns
```

```
['longitude',
 'latitude',
 'housing_median_age',
 'total_rooms',
 'total_bedrooms',
 'population',
 'households',
 'median_income',
 'median_house_value']
```

```
len(numeric_columns) , len(data.columns)
```

```
(9, 10)
```

```
numeric_columns.remove('longitude')
numeric_columns.remove('latitude')
```

```
temp_data = data[numeric_columns]
temp_data
```

	housing_median_age	total_rooms	total_bedrooms	population	households	medv
0	41.0	880.0	129.0	322.0	126.0	
1	21.0	7099.0	1106.0	2401.0	1138.0	
2	52.0	1467.0	190.0	496.0	177.0	
3	52.0	1274.0	235.0	558.0	219.0	
4	52.0	1627.0	280.0	565.0	259.0	
...	...	...	...	...	...	...
20635	25.0	1665.0	374.0	845.0	330.0	
20636	18.0	697.0	150.0	356.0	114.0	
20637	17.0	2254.0	485.0	1007.0	433.0	
20638	18.0	1860.0	409.0	741.0	349.0	

## Normalization

20640 rows x 7 columns

```
from sklearn.preprocessing import StandardScaler , MinMaxScaler
import warnings
warnings.filterwarnings('ignore')
normalizer = MinMaxScaler()
temp_data.dropna(axis = 1 , inplace = True)
normalized_data = normalizer.fit_transform(temp_data)
pd.DataFrame(normalized_data , columns = temp_data.columns)
```

	housing_median_age	total_rooms	population	households	median_income	medv
0	0.784314	0.022331	0.008941	0.020556	0.539668	
1	0.392157	0.180503	0.067210	0.186976	0.538027	
2	1.000000	0.037260	0.013818	0.028943	0.466028	
3	1.000000	0.032352	0.015555	0.035849	0.354699	
4	1.000000	0.041330	0.015752	0.042427	0.230776	
...	...	...	...	...	...	...
20635	0.470588	0.042296	0.023599	0.054103	0.073130	
20636	0.333333	0.017676	0.009894	0.018582	0.141853	
20637	0.313725	0.057277	0.028140	0.071041	0.082764	
20638	0.333333	0.047256	0.020684	0.057227	0.094295	
20639	0.294118	0.070782	0.038790	0.086992	0.130253	

20640 rows x 6 columns

## Standardization

```

standard_scaler = StandardScaler()
standardized_data = standard_scaler.fit_transform(temp_data)
pd.DataFrame(standardized_data , columns = temp_data.columns)

```

	housing_median_age	total_rooms	population	households	median_income	mec
<b>0</b>	0.982143	-0.804819	-0.974429	-0.977033	2.344766	
<b>1</b>	-0.607019	2.045890	0.861439	1.669961	2.332238	
<b>2</b>	1.856182	-0.535746	-0.820777	-0.843637	1.782699	
<b>3</b>	1.856182	-0.624215	-0.766028	-0.733781	0.932968	
<b>4</b>	1.856182	-0.462404	-0.759847	-0.629157	-0.012881	
...	...	...	...	...	...	
<b>20635</b>	-0.289187	-0.444985	-0.512592	-0.443449	-1.216128	
<b>20636</b>	-0.845393	-0.888704	-0.944405	-1.008420	-0.691593	
<b>20637</b>	-0.924851	-0.174995	-0.369537	-0.174042	-1.142593	
<b>20638</b>	-0.845393	-0.355600	-0.604429	-0.393753	-1.054583	
<b>20639</b>	-1.004309	0.068408	-0.033977	0.079672	-0.780129	

20640 rows × 6 columns

## Handling With Missing Values

```
data.isnull().sum()
```

```

longitude          0
latitude           0
housing_median_age  0
total_rooms         0
total_bedrooms     207
population          0
households          0
median_income       0
median_house_value  0
ocean_proximity     0
dtype: int64

```

```

# here I Will show you imputing values in Null columns only for 'agent' column
data['total_bedrooms'].isnull().sum()

```

207

## Simple Imputer

```

from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values=np.nan , strategy='mean')
agent_col = imputer.fit_transform(data['total_bedrooms'].values.reshape(-1,1))
pd.DataFrame(agent_col).isnull().sum()

```

```
pd.DataFrame(agent_col).isnull().sum()
```

```
0    0
dtype: int64
```

```
data['total_bedrooms'].isnull().sum()
```

```
207
```

## Discretization

```
from sklearn.preprocessing import KBinsDiscretizer
temp_data.head()
```

	housing_median_age	total_rooms	population	households	median_income	median_hoi
<b>0</b>	41.0	880.0	322.0	126.0	8.3252	
<b>1</b>	21.0	7099.0	2401.0	1138.0	8.3014	
<b>2</b>	52.0	1467.0	496.0	177.0	7.2574	
<b>3</b>	52.0	1274.0	558.0	219.0	5.6431	
<b>4</b>	52.0	1627.0	565.0	259.0	3.8462	

## Quantile Discretization Transform

```
trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='quantile')
new_data = trans.fit_transform(temp_data)
pd.DataFrame(new_data,columns = temp_data.columns )
```

## Uniform Discretization Transform

```
trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='uniform')
new_data = trans.fit_transform(temp_data)
```

```
pd.DataFrame(new_data,columns = temp_data.columns )
```

	housing_median_age	total_rooms	population	households	median_income	medv
<b>0</b>	7.0	0.0	0.0	0.0	5.0	
<b>1</b>	3.0	1.0	0.0	1.0	5.0	
<b>2</b>	9.0	0.0	0.0	0.0	4.0	
<b>3</b>	9.0	0.0	0.0	0.0	3.0	
<b>4</b>	9.0	0.0	0.0	0.0	2.0	
...	...	...	...	...	...	
<b>20635</b>	4.0	0.0	0.0	0.0	0.0	
<b>20636</b>	3.0	0.0	0.0	0.0	1.0	
<b>20637</b>	3.0	0.0	0.0	0.0	0.0	
<b>20638</b>	3.0	0.0	0.0	0.0	0.0	
<b>20639</b>	2.0	0.0	0.0	0.0	1.0	

20640 rows × 6 columns

## KMeans Discretization Transform

```
trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='kmeans')
new_data = trans.fit_transform(temp_data)
```

```
pd.DataFrame(new_data,columns = temp_data.columns )
```

	housing_median_age	total_rooms	population	households	median_income	mec
0	7.0	0.0	0.0	0.0	6.0	
1	3.0	4.0	3.0	4.0	6.0	
2	9.0	1.0	0.0	0.0	6.0	
3	9.0	0.0	0.0	0.0	4.0	
4	9.0	1.0	0.0	1.0	3.0	
...	...	...	...	...	...	
20635	4.0	1.0	1.0	1.0	0.0	
20636	3.0	0.0	0.0	0.0	1.0	
20637	3.0	1.0	1.0	2.0	0.0	
20638	3.0	1.0	0.0	1.0	0.0	
20639	2.0	2.0	1.0	2.0	1.0	

20640 rows × 6 columns